

fault_detection

March 12, 2019

```
In [3]: import pandas as pd
import pyarrow.parquet as pq
import os
import pyarrow
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from scipy.signal import *
import statsmodels.api as sm
from keras import backend as K
import gc
from sklearn.feature_selection import f_classif
import lightgbm as lgbm
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import expon, uniform, norm
from scipy.stats import randint, poisson
from sklearn.metrics import confusion_matrix, make_scorer

print(os.listdir('../input'))
```

```
['metadata_train.csv', 'train.parquet', 'train_subset_meta.csv', 'metadata_train_V2.csv', 'subse
```

1 Import data

```
In [4]: %%time

print('load data...')
train_df = pq.read_pandas("../input/train.parquet").to_pandas()
#test_df = pq.read_pandas("../input/test.parquet").to_pandas()
train_meta_df = pd.read_csv("../input/metadata_train.csv")
#test_meta_df = pd.read_csv("../input/metadata_test.csv")
print('import ok')
```

load data...

import ok

CPU times: user 55.6 s, sys: 11.6 s, total: 1min 7s

Wall time: 1min 8s

```
In [5]: train_meta_df.head(n=9)
```

```
Out[5]:
```

	signal_id	id_measurement	phase	target
0	0	0	0	0
1	1	0	1	0
2	2	0	2	0
3	3	1	0	1
4	4	1	1	1
5	5	1	2	1
6	6	2	0	0
7	7	2	1	0
8	8	2	2	0

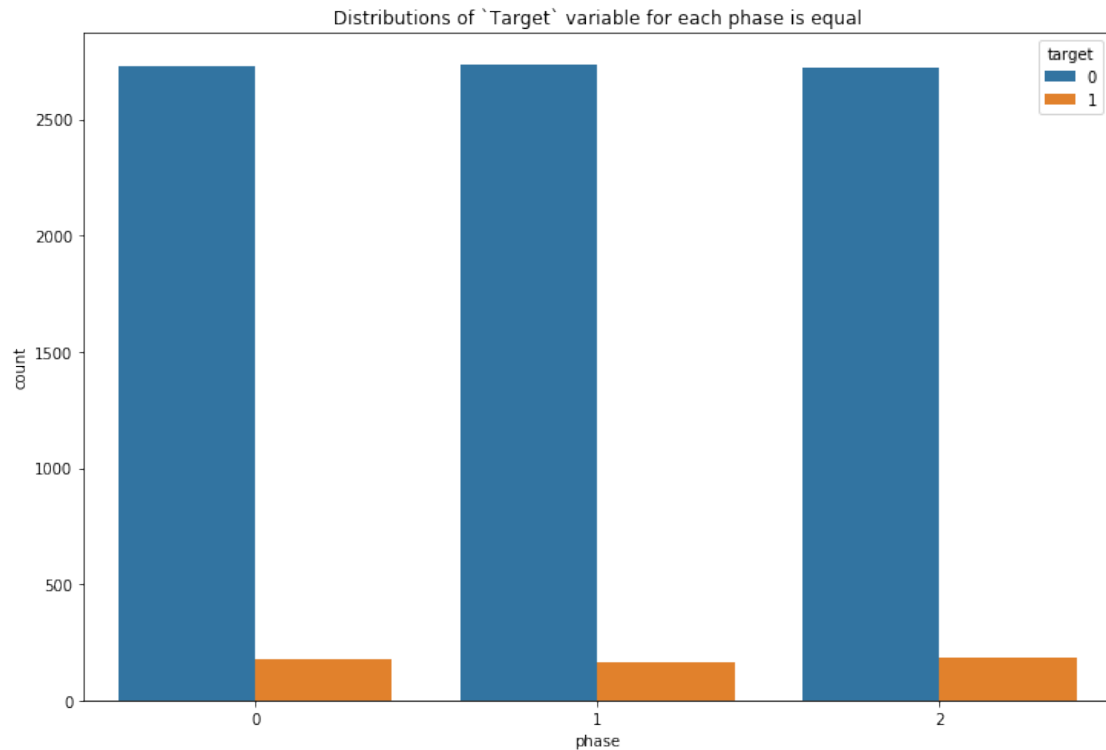
2 Exploratory data analysis

```
In [6]: # plot settings
```

```
rand_seed = 135  
np.random.seed(rand_seed)  
xsize = 12.0  
ysize = 8.0
```


```
In [7]: fig, ax = plt.subplots()  
fig.set_size_inches(xsize, ysize)
```

```
ax = sns.countplot(x="phase", hue="target", data=train_meta_df, ax=ax)  
ax.set_title("Distributions of `Target` variable for each phase is equal")  
plt.show()
```



```
In [8]: %%time
```

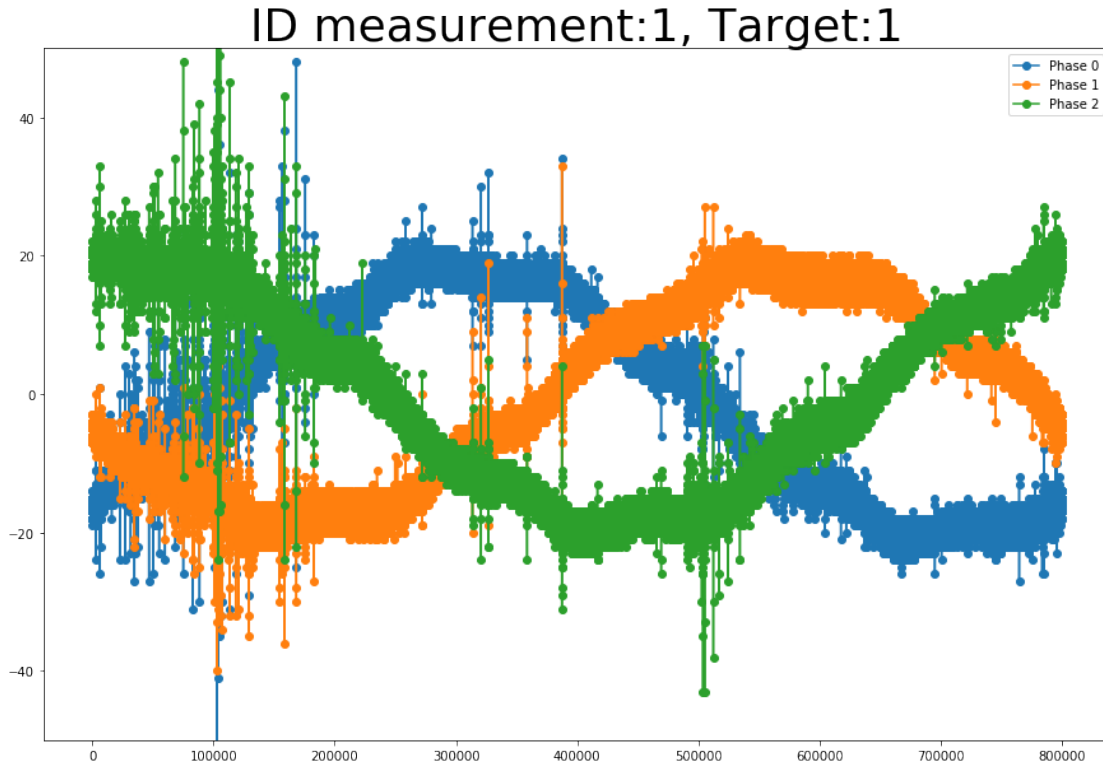
```
plt.figure(figsize=(15, 10))
plt.title("ID measurement:0, Target:0", fontdict={'fontsize':36})
plt.plot(train_df["0"].values, marker="o", label='Phase 0')
plt.plot(train_df["1"].values, marker="o", label='Phase 1')
plt.plot(train_df["2"].values, marker="o", label='Phase 2')
plt.ylim(-50,50)
plt.legend()
plt.show()
```



CPU times: user 18.3 s, sys: 364 ms, total: 18.6 s
 Wall time: 10 s

In [9]: %%time

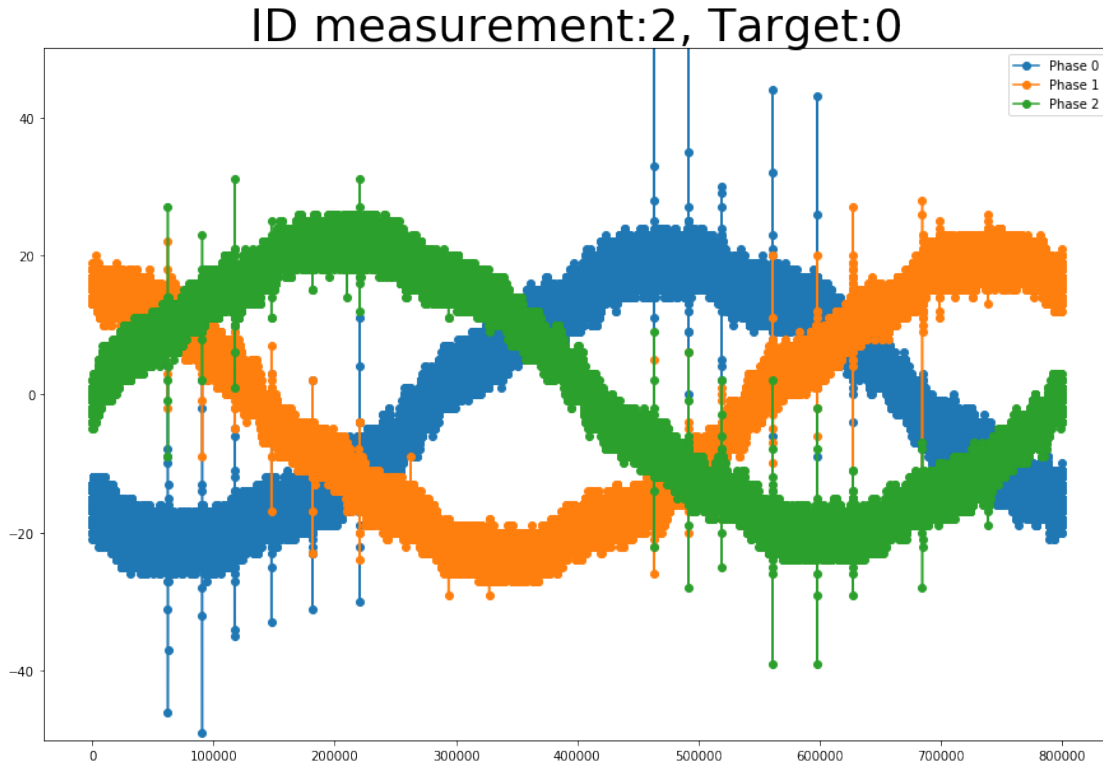
```
plt.figure(figsize=(15, 10))
plt.title("ID measurement:1, Target:1", fontdict={'fontsize':36})
plt.plot(train_df["3"].values, marker="o", label='Phase 0')
plt.plot(train_df["4"].values, marker="o", label='Phase 1')
plt.plot(train_df["5"].values, marker="o", label='Phase 2')
plt.ylim(-50,50)
plt.legend()
plt.show()
```



CPU times: user 16.2 s, sys: 364 ms, total: 16.5 s
 Wall time: 8.88 s

In [10]: %%time

```
plt.figure(figsize=(15, 10))
plt.title("ID measurement:2, Target:0", fontdict={'fontsize':36})
plt.plot(train_df["6"].values, marker="o", label='Phase 0')
plt.plot(train_df["7"].values, marker="o", label='Phase 1')
plt.plot(train_df["8"].values, marker="o", label='Phase 2')
plt.ylim(-50,50)
plt.legend()
plt.show()
```



CPU times: user 16.7 s, sys: 332 ms, total: 17 s
 Wall time: 8.95 s

3 Save a subset of data

In order to speed up the code, we only run on a subset of the training data.

```
In [11]: train_subset_df = train_df.iloc[:,range(0,99)]
         train_subset_meta_df = train_meta_df.iloc[range(0,99),:]
         #train_subset_df.to_csv('../input/train_subset.csv')
         #train_subset_meta_df.to_csv('../input/train_subset_meta.csv')

         # uncomment to use the full dataset after
         train_subset_df = train_df
         train_subset_meta_df = train_meta_df
```

```
In [12]: train_subset_meta_df.head(n=9)
```

```
Out[12]:
```

	signal_id	id_measurement	phase	target
0	0	0	0	0
1	1	0	1	0
2	2	0	2	0

3	3	1	0	1
4	4	1	1	1
5	5	1	2	1
6	6	2	0	0
7	7	2	1	0
8	8	2	2	0

4 Feature engineering

4.1 Mean, median and standard deviation

In [13]: %%time

```
mean_list = train_subset_df.apply(np.mean)
median_list = train_subset_df.apply(np.median)
std_list = train_subset_df.apply(np.std)
```

CPU times: user 3min 38s, sys: 0 ns, total: 3min 38s

Wall time: 1min 3s

```
In [14]: mean_signal_df = mean_list.to_frame()
mean_signal_df = mean_signal_df.reset_index()
mean_signal_df = mean_signal_df.drop("index",axis=1)

train_subset_meta_df = train_subset_meta_df.merge(mean_signal_df,"inner", left_index=True, right_index=True)
train_subset_meta_df = train_subset_meta_df.rename(index=str, columns={0:"mean"})

median_signal_df = median_list.to_frame()

train_subset_meta_df = train_subset_meta_df.merge(median_signal_df,"inner", left_index=True, right_index=True)
train_subset_meta_df = train_subset_meta_df.rename(index=str, columns={0:"median"})

std_signal_df = std_list.to_frame()

train_subset_meta_df = train_subset_meta_df.merge(std_signal_df,"inner", left_index=True, right_index=True)
train_subset_meta_df = train_subset_meta_df.rename(index=str, columns={0:"std_dev"})
```


In [15]: train_subset_meta_df.head(n=9)

```
Out[15]:
```


	signal_id	id_measurement	phase	target	mean	median	std_dev
0	0	0	0	0	-0.960271	-1.0	13.870724
1	1	0	1	0	-0.194125	0.0	13.037134
2	2	0	2	0	-0.043555	0.0	13.684282
3	3	1	0	1	-0.997401	-1.0	13.673630
4	4	1	1	1	-0.175586	0.0	12.938372
5	5	1	2	1	-0.036004	0.0	13.545777
6	6	2	0	0	-1.146185	-1.0	14.064211

7	7	2	1	0	-1.952695	-2.0	14.774424
8	8	2	2	0	0.873370	1.0	14.815668

```
In [16]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Mean Across Target Variables")
ax = sns.boxplot(x="target", y="mean", data=train_subset_meta_df)
```





```
In [17]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Median Across Target Variables")
ax = sns.boxplot(x="target", y="median", data=train_subset_meta_df)
```



```
In [18]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
```

```
plt.title("Standard Deviation Across Target Variables")  
ax = sns.boxplot(x="target", y="std_dev", data=train_subset_meta_df)
```

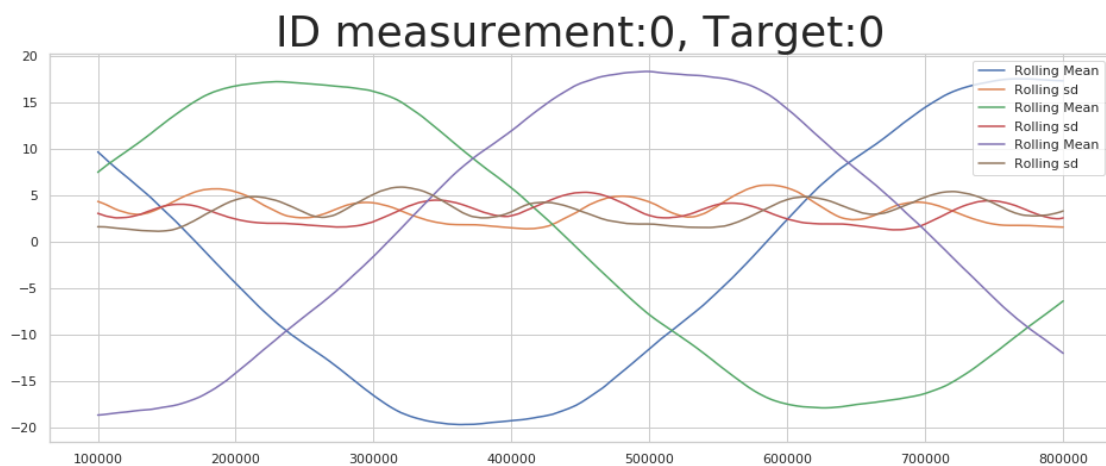


4.2 Amplitude of Rolling Series

Let's smooth the time series.


```
In [19]: ts1 = train_df["0"]
         ts2 = train_df["1"]
         ts3 = train_df["2"]
```

```
plt.figure(figsize=(16,6))
plt.title("ID measurement:0, Target:0", fontdict={'fontsize':36})
plt.plot(ts1.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts1.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.plot(ts2.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts2.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.plot(ts3.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts3.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.legend();
```




```
In [20]: ts1 = train_df["3"]
         ts2 = train_df["4"]
         ts3 = train_df["5"]
```

```
plt.figure(figsize=(16,6))
plt.title("ID measurement:1, Target:1", fontdict={'fontsize':36})
plt.plot(ts1.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts1.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.plot(ts2.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts2.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.plot(ts3.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts3.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.legend();
```



```
In [21]: ts1 = train_df["6"]
        ts2 = train_df["7"]
        ts3 = train_df["8"]
```

```
plt.figure(figsize=(16,6))
plt.title("ID measurement:2, Target:0", fontdict={'fontsize':36})
plt.plot(ts1.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts1.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.plot(ts2.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts2.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.plot(ts3.rolling(window=100000,center=False).mean(),label='Rolling Mean');
plt.plot(ts3.rolling(window=100000,center=False).std(),label='Rolling sd');
plt.legend();
```



Let's look at the amplitude across each target group. To calculate the amplitude, I smooth the powerline signals to create a single wave then I subtract the lowest and highest point.

```

In [22]: def calc_rolling_amp(row, window=100000):
          return np.max(row.rolling(window,center=False).mean()) - np.min(row.rolling(window=

rolling100k_amp = train_subset_df.apply(calc_rolling_amp)

In [23]: rolling100k_amp_df = rolling100k_amp.to_frame()
train_subset_meta_df = train_subset_meta_df.merge(rolling100k_amp_df,"inner", left_index=True, right_index=True)
train_subset_meta_df = train_subset_meta_df.rename(index=str, columns={0:"rolling100k_amp"})

In [24]: train_subset_meta_df.head(n=9)

```

Out[24]:


	signal_id	id_measurement	phase	target	mean	median	std_dev	\
0	0	0	0	0	-0.960271	-1.0	13.870724	
1	1	0	1	0	-0.194125	0.0	13.037134	
2	2	0	2	0	-0.043555	0.0	13.684282	
3	3	1	0	1	-0.997401	-1.0	13.673630	
4	4	1	1	1	-0.175586	0.0	12.938372	
5	5	1	2	1	-0.036004	0.0	13.545777	
6	6	2	0	0	-1.146185	-1.0	14.064211	
7	7	2	1	0	-1.952695	-2.0	14.774424	
8	8	2	2	0	0.873370	1.0	14.815668	

	rolling100k_amp
0	37.21537
1	35.10791
2	36.97624
3	37.53126
4	35.35856
5	36.87904
6	39.47469
7	41.58219
8	41.58396

```

In [25]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Amplitude Across Target Variables")
ax = sns.boxplot(x="target", y="rolling100k_amp", data=train_subset_meta_df)

```



4.3 Measuring Amount of Noisy Points

4.3.1 Number of points 1SD from the mean

```
In [26]: def count1SDfromTheMean(row):  
          max_1sd = np.mean(row) + np.std(row)  
          min_1sd = np.mean(row) - np.std(row)
```

```
noise_points = [x for x in row if (x > max_1sd) or (x < min_1sd)]
return (len(noise_points))
```

```
In [27]: count1SDfromTheMean_list = train_subset_df.apply(count1SDfromTheMean)
count1SDfromTheMean_df = count1SDfromTheMean_list.to_frame()
train_subset_meta_df = train_subset_meta_df.merge(count1SDfromTheMean_df,"inner", left_
train_subset_meta_df = train_subset_meta_df.rename(index=str, columns={0:"count1SDfromT


In [28]: train_subset_meta_df.head(n=9)
```

```
Out[28]:
```

	signal_id	id_measurement	phase	target	mean	median	std_dev	\
0	0	0	0	0	-0.960271	-1.0	13.870724	
1	1	0	1	0	-0.194125	0.0	13.037134	
2	2	0	2	0	-0.043555	0.0	13.684282	
3	3	1	0	1	-0.997401	-1.0	13.673630	
4	4	1	1	1	-0.175586	0.0	12.938372	
5	5	1	2	1	-0.036004	0.0	13.545777	
6	6	2	0	0	-1.146185	-1.0	14.064211	
7	7	2	1	0	-1.952695	-2.0	14.774424	
8	8	2	2	0	0.873370	1.0	14.815668	

	rolling100k_amp	count1SDfromTheMean
0	37.21537	377353
1	35.10791	372859
2	36.97624	377776
3	37.53126	381716
4	35.35856	377552
5	36.87904	379631
6	39.47469	378400
7	41.58219	391669
8	41.58396	394043

```
In [29]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Noise Count Across Target Variables")
ax = sns.boxplot(x="target", y="count1SDfromTheMean", data=train_subset_meta_df)
```



4.3.2 Number of points 2SD from the mean

```
In [30]: def count2SDfromTheMean(row):
max_1sd = np.mean(row) + (2 * np.std(row))
min_1sd = np.mean(row) - (2 * np.std(row))
noise_points = [x for x in row if (x > max_1sd) or (x < min_1sd)]
return (len(noise_points))
```



```

In [31]: count2SDfromTheMean_list = train_subset_df.apply(count2SDfromTheMean)
count2SDfromTheMean_df = count2SDfromTheMean_list.to_frame()
train_subset_meta_df = train_subset_meta_df.merge(count2SDfromTheMean_df,"inner", left_
train_subset_meta_df = train_subset_meta_df.rename(index=str, columns={0:"count2SDfromT

In [32]: train_subset_meta_df.head()

Out[32]:
```

	signal_id	id_measurement	phase	target	mean	median	std_dev	\
0	0	0	0	0	-0.960271	-1.0	13.870724	
1	1	0	1	0	-0.194125	0.0	13.037134	
2	2	0	2	0	-0.043555	0.0	13.684282	
3	3	1	0	1	-0.997401	-1.0	13.673630	
4	4	1	1	1	-0.175586	0.0	12.938372	


	rolling100k_amp	count1SDfromTheMean	count2SDfromTheMean
0	37.21537	377353	21
1	35.10791	372859	7
2	36.97624	377776	23
3	37.53126	381716	28
4	35.35856	377552	24


```

In [33]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Noise Count Across Target Variables")
ax = sns.boxplot(x="target", y="count2SDfromTheMean", data=train_subset_meta_df)
plt.ylim(0,250)

Out[33]: (0, 250)

```




4.4 FFT


```
In [34]: plt.figure(figsize=(15, 10))
plt.title("ID measurement:2, Target:0", fontdict={'fontsize':36})
plt.plot(np.fft.fft(train_subset_df["6"]), marker="o", label='Phase 0')
plt.plot(np.fft.fft(train_subset_df["7"]), marker="o", label='Phase 1')
plt.plot(np.fft.fft(train_subset_df["8"]), marker="o", label='Phase 2')
```

```
plt.legend()
plt.show()
```

```
/home/pierre/bin/anaconda3/lib/python3.6/site-packages/numpy/core/numeric.py:538: ComplexWarning
return array(a, dtype, copy=False, order=order)
```



```
In [35]: plt.figure(figsize=(15, 10))
plt.title("ID measurement:1, Target:1", fontdict={'fontsize':36})
plt.plot(np.fft.fft(train_subset_df["3"]), marker="o", label='Phase 0')
plt.plot(np.fft.fft(train_subset_df["4"]), marker="o", label='Phase 1')
plt.plot(np.fft.fft(train_subset_df["5"]), marker="o", label='Phase 2')
plt.legend()
plt.show()
```



4.5 Power Spectral density using the Welch's method

```
In [36]: fig, axes = plt.subplots(nrows=2, ncols=2)
fig.set_size_inches(2.0*xsize, 2.0*ysize)
axes = axes.flatten()

f, Pxx = welch(train_subset_df["0"].values)
axes[0].plot(f, Pxx, marker="o", linestyle="none")
axes[0].set_title("Signal ID: 0")
axes[0].axhline(y=2.5, color="k", linestyle="--")

f, Pxx = welch(train_subset_df["1"].values)
axes[1].plot(f, Pxx, marker="o", linestyle="none")
axes[1].set_title("Signal ID: 1")
axes[1].axhline(y=2.5, color="k", linestyle="--")

f, Pxx = welch(train_subset_df["2"].values)
axes[2].plot(f, Pxx, marker="o", linestyle="none")
axes[2].set_title("Signal ID: 2")
axes[2].axhline(y=2.5, color="k", linestyle="--")

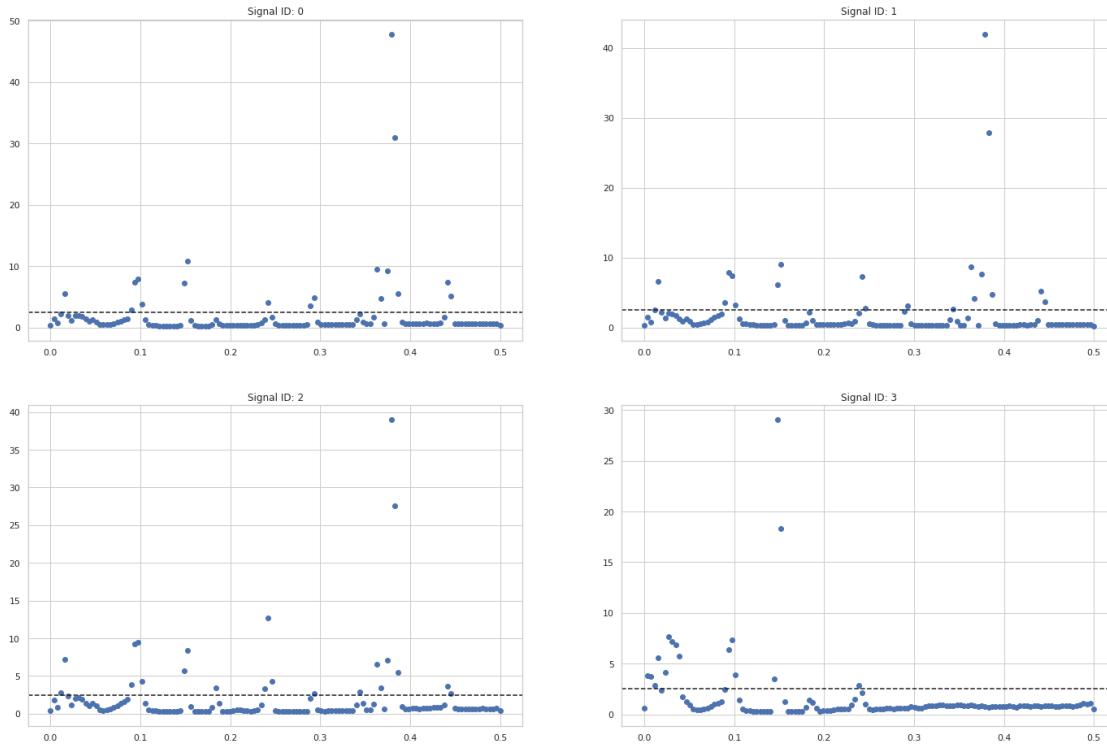
f, Pxx = welch(train_subset_df["3"].values)
axes[3].plot(f, Pxx, marker="o", linestyle="none")
```

```

axes[3].set_title("Signal ID: 3")
axes[3].axhline(y=2.5, color="k", linestyle="--")

plt.show()

```



```
In [37]: %%time
```

```

def welch_max_power_and_frequency(signal):
    f, Pxx = welch(signal)
    ix = np.argmax(Pxx)
    strong_count = np.sum(Pxx>2.5)
    avg_amp = np.mean(Pxx)
    sum_amp = np.sum(Pxx)
    std_amp = np.std(Pxx)
    median_amp = np.median(Pxx)
    return [Pxx[ix], f[ix], strong_count, avg_amp, sum_amp, std_amp, median_amp]

```

```
power_spectrum_summary = train_subset_df.apply(welch_max_power_and_frequency, result_ty
```

```
CPU times: user 26min 30s, sys: 24 ms, total: 26min 30s
```

```
Wall time: 4min 41s
```

```
In [38]: power_spectrum_summary = power_spectrum_summary.T.rename(columns={0:"max_amp", 1:"max_f",
4:"sum_amp", 5:"std_a",
power_spectrum_summary.head()
```

```
Out [38]:
```

	max_amp	max_freq	strong_amp_count	avg_amp	sum_amp	std_amp	\
0	47.831062	0.378906	18.0	1.995473	257.415955	5.220994	
1	41.982578	0.378906	20.0	1.836459	236.903198	4.641738	
2	38.999954	0.378906	22.0	1.935738	249.710236	4.503104	
3	29.060942	0.148438	16.0	1.590296	205.148132	3.224866	
4	20.080660	0.148438	19.0	1.375932	177.495178	2.577763	

	median_amp
0	0.612473
1	0.458779
2	0.652928
3	0.798159
4	0.439266

```
In [39]: power_spectrum_summary.index = power_spectrum_summary.index.astype(int)

train_subset_meta_df = train_subset_meta_df.merge(power_spectrum_summary, left_on="signal_id", right_on="id_measurement")
train_subset_meta_df.head()
```


```
Out [39]:
```

	signal_id	id_measurement	phase	target	mean	median	std_dev	\
0	0	0	0	0	-0.960271	-1.0	13.870724	
1	1	0	1	0	-0.194125	0.0	13.037134	
2	2	0	2	0	-0.043555	0.0	13.684282	
3	3	1	0	1	-0.997401	-1.0	13.673630	
4	4	1	1	1	-0.175586	0.0	12.938372	


	rolling100k_amp	count1SDfromTheMean	count2SDfromTheMean	max_amp	\
0	37.21537	377353	21	47.831062	
1	35.10791	372859	7	41.982578	
2	36.97624	377776	23	38.999954	
3	37.53126	381716	28	29.060942	
4	35.35856	377552	24	20.080660	

	max_freq	strong_amp_count	avg_amp	sum_amp	std_amp	median_amp
0	0.378906	18.0	1.995473	257.415955	5.220994	0.612473
1	0.378906	20.0	1.836459	236.903198	4.641738	0.458779
2	0.378906	22.0	1.935738	249.710236	4.503104	0.652928
3	0.148438	16.0	1.590296	205.148132	3.224866	0.798159
4	0.148438	19.0	1.375932	177.495178	2.577763	0.439266


```
In [40]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Max amplitude PSD")
ax = sns.boxplot(x="target", y="max_amp", data=train_subset_meta_df)
```




```
In [41]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Max frequency PSD")
ax = sns.boxplot(x="target", y="max_freq", data=train_subset_meta_df)
```




```
In [42]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Strong amplitude count")
ax = sns.boxplot(x="target", y="strong_amp_count", data=train_subset_meta_df)
```


```
In [43]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Average amplitude")
ax = sns.boxplot(x="target", y="avg_amp", data=train_subset_meta_df)
```




```
In [44]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Sum amp")
ax = sns.boxplot(x="target", y="sum_amp", data=train_subset_meta_df)
```



```
In [45]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("Std amp")
ax = sns.boxplot(x="target", y="std_amp", data=train_subset_meta_df)
```



```
In [46]: plt.figure(figsize=(6,8))
sns.set(style="whitegrid")
plt.title("median amp")
ax = sns.boxplot(x="target", y="median_amp", data=train_subset_meta_df)
```



5 Find the most important features

```
In [47]: X_cols = ["phase"] + train_subset_meta_df.columns[4:].tolist()  
         X_cols
```

```
Out[47]: ['phase',  
          'mean',
```

```

'median',
'std_dev',
'rolling100k_amp',
'count1SDfromTheMean',
'count2SDfromTheMean',
'max_amp',
'max_freq',
'strong_amp_count',
'avg_amp',
'sum_amp',
'std_amp',
'median_amp']

```

```
In [48]: Fvals, pvals = f_classif(train_subset_meta_df[X_cols], train_subset_meta_df["target"])
```

```

print("F-value | P-value | Feature Name")
print("-----")

for i, col in enumerate(X_cols):
    print("%.4f"%Fvals[i]+" | "+"%.4f"%pvals[i]+" | "+col)

```

```

F-value | P-value | Feature Name
-----
0.0274 | 0.8686 | phase
9.2846 | 0.0023 | mean
7.0488 | 0.0079 | median
6.3229 | 0.0119 | std_dev
25.2459 | 0.0000 | rolling100k_amp
8.1104 | 0.0044 | count1SDfromTheMean
16.1775 | 0.0001 | count2SDfromTheMean
5.3612 | 0.0206 | max_amp
14.1640 | 0.0002 | max_freq
112.3396 | 0.0000 | strong_amp_count
62.2001 | 0.0000 | avg_amp
62.2001 | 0.0000 | sum_amp
3.9723 | 0.0463 | std_amp
250.2839 | 0.0000 | median_amp

```

So as expected phase is a useless feature on its own, but interestingly std_amp, median_amp, signal_std, max_amp may not be extremely useful variables because we cannot reject the null with a significance of 0.01 for these. However the features signal_mean, signal_sum, max_freq, strong_amp_count, avg_amp, and sum_amp all look like very useful features, even on their own.

```
In [49]: train_subset_meta_df.to_csv('../input/metadata_train_V2.csv')
         train_subset_meta_df.head(n=9)
```

```

Out[49]:   signal_id  id_measurement  phase  target      mean  median  std_dev \
         0         0         0         0         0 -0.960271   -1.0  13.870724

```

1	1	0	1	0	-0.194125	0.0	13.037134
2	2	0	2	0	-0.043555	0.0	13.684282
3	3	1	0	1	-0.997401	-1.0	13.673630
4	4	1	1	1	-0.175586	0.0	12.938372
5	5	1	2	1	-0.036004	0.0	13.545777
6	6	2	0	0	-1.146185	-1.0	14.064211
7	7	2	1	0	-1.952695	-2.0	14.774424
8	8	2	2	0	0.873370	1.0	14.815668

	rolling100k_amp	count1SDfromTheMean	count2SDfromTheMean	max_amp \
0	37.21537	377353	21	47.831062
1	35.10791	372859	7	41.982578
2	36.97624	377776	23	38.999954
3	37.53126	381716	28	29.060942
4	35.35856	377552	24	20.080660
5	36.87904	379631	68	23.466799
6	39.47469	378400	22	120.421745
7	41.58219	391669	1	81.411369
8	41.58396	394043	6	123.972298

	max_freq	strong_amp_count	avg_amp	sum_amp	std_amp	median_amp
0	0.378906	18.0	1.995473	257.415955	5.220994	0.612473
1	0.378906	20.0	1.836459	236.903198	4.641738	0.458779
2	0.378906	22.0	1.935738	249.710236	4.503104	0.652928
3	0.148438	16.0	1.590296	205.148132	3.224866	0.798159
4	0.148438	19.0	1.375932	177.495178	2.577763	0.439266
5	0.148438	19.0	1.653171	213.259018	2.911619	0.874973
6	0.382812	17.0	4.404845	568.224976	15.451883	0.395055
7	0.382812	14.0	3.057476	394.414368	10.961865	0.334526
8	0.382812	16.0	3.072290	396.325439	12.887474	0.324272

6 Fit libGBM model and hyperparameter tuning

```
In [52]: from keras import backend as K
         from sklearn.metrics import matthews_corrcoef
```

```
In [54]: #def mcc(y_true, y_pred, labels=None, sample_weight=None):
         #     """ Matthew's coefficient correlation """
         #     tn, fp, fn, tp = confusion_matrix(y_true, y_pred, labels=labels, sample_weight=sam
         #     mcc = (tp*tn - fp*fn)/np.sqrt((tp + fp)*(tp + fn)*(tn + fp)*(tn + fn))
         #     return mcc

         #def mcc(y_true, y_pred):
         #     '''Calculates the Matthews correlation coefficient measure for quality
         #     of binary classification problems.
         #     '''
```

```

#     y_pred_pos = K.round(K.clip(y_pred, 0, 1))
#     y_pred_neg = 1 - y_pred_pos
#
#     y_pos = K.round(K.clip(y_true, 0, 1))
#     y_neg = 1 - y_pos
#
#     tp = K.sum(y_pos * y_pred_pos)
#     tn = K.sum(y_neg * y_pred_neg)
#
#     fp = K.sum(y_neg * y_pred_pos)
#     fn = K.sum(y_pos * y_pred_neg)
#
#     numerator = (tp * tn - fp * fn)
#     denominator = K.sqrt((tp + fp) * (tp + fn) * (tn + fp) * (tn + fn))
#
#     return numerator / (denominator + K.epsilon())

def mcc(y_true, y_pred):
    return matthews_corrcoef(y_true, y_pred)

mcc_scorer = make_scorer(mcc)

lgbm_classifier = lgbm.LGBMClassifier(boosting_type='gbdt', max_depth=-1, subsample_for
                                     class_weight=None, min_split_gain=0.0, min_child_
                                     subsample_freq=0, random_state=rand_seed, n_jobs=

param_distributions = {
    "num_leaves": randint(16, 48),
    "learning_rate": expon(),
    "reg_alpha": expon(),
    "reg_lambda": expon(),
    "colsample_bytree": uniform(0.25, 1.0),
    "min_child_samples": randint(10, 30),
    "n_estimators": randint(50, 250)
}

clf = RandomizedSearchCV(lgbm_classifier, param_distributions, n_iter=100, scoring=mcc_
                        refit=True, cv=5, verbose=1, random_state=rand_seed, error_sco

clf.fit(train_subset_meta_df[X_cols], train_subset_meta_df["target"])

```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
/home/pierre/bin/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_validation.py:55
lightgbm.basic.LightGBMError: Check failed: feature_fraction <=1.0 at /tmp/pip-req-build-ztyh0k2


```

FitFailedWarning)
/home/pierre/bin/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:543: Ru
    mcc = cov_ytyp / np.sqrt(cov_ytyt * cov_ypyp)
[Parallel(n_jobs=1)]: Done 500 out of 500 | elapsed: 3.0min finished

```

```

Out [54]: RandomizedSearchCV(cv=5, error_score=-1.0,
    estimator=LGBMClassifier(boosting_type='gbdt', class_weight=None, colsample_b
    importance_type='split', learning_rate=0.1, max_depth=-1,
    min_child_samples=20, min_child_weight=0.001, min_split_gain=0.0,
    n_estimators=100, n_jobs=1, num_leaves=31, objective='binary',
    random_state=135, reg_alpha=0.0, reg_lambda=0.0, silent=True,
    subsample=1.0, subsample_for_bin=200000, subsample_freq=0),
    fit_params=None, iid=True, n_iter=100, n_jobs=1,
    param_distributions={'num_leaves': <scipy.stats._distn_infrastructure.rv_froz
    pre_dispatch='2*n_jobs', random_state=135, refit=True,
    return_train_score=True, scoring=make_scorer(mcc), verbose=1)

```

```

In [55]: print(clf.best_score_)

```

```

0.5401876970150516

```

```

In [57]: clf.best_estimator_

```

```

Out [57]: LGBMClassifier(boosting_type='gbdt', class_weight=None,
    colsample_bytree=0.456616770229323, importance_type='split',
    learning_rate=0.4118819089418852, max_depth=-1,
    min_child_samples=24, min_child_weight=0.001, min_split_gain=0.0,
    n_estimators=197, n_jobs=1, num_leaves=19, objective='binary',
    random_state=135, reg_alpha=0.28274323494050946,
    reg_lambda=1.5600109217504974, silent=True, subsample=1.0,
    subsample_for_bin=200000, subsample_freq=0)

```

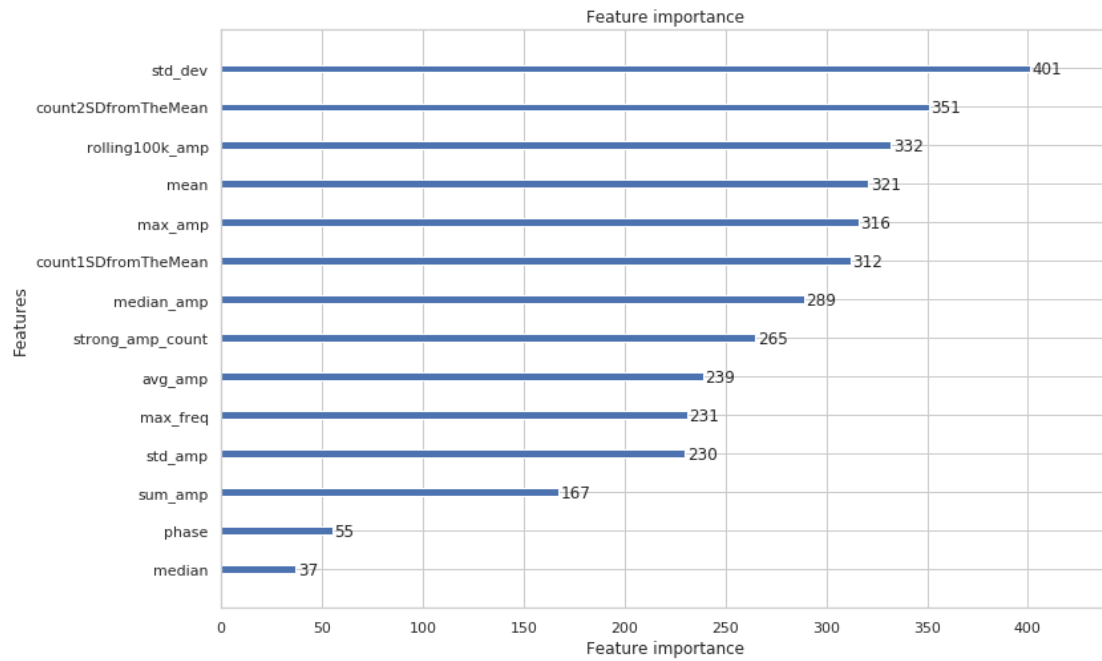
```

In [58]: fig, ax = plt.subplots()
    fig.set_size_inches(xsize, ysize)

    lgbm.plot_importance(clf.best_estimator_, ax=ax)

    plt.show()

```



In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: